

Parameterized Bounded-Depth Frege is Not Optimal

Olaf Beyersdorff^{1*}, Nicola Galesi^{2**}, Massimo Lauria², and
Alexander Razborov^{3***}

¹ Institut für Theoretische Informatik, Leibniz Universität Hannover, Germany

² Dipartimento di Informatica, Sapienza Università di Roma, Italy

³ Department of Computer Science, The University of Chicago

Abstract. A general framework for parameterized proof complexity was introduced by Dantchev, Martin, and Szeider [9]. There the authors concentrate on tree-like Parameterized Resolution—a parameterized version of classical Resolution—and their gap complexity theorem implies lower bounds for that system.

The main result of the present paper significantly improves upon this by showing optimal lower bounds for a parameterized version of bounded-depth Frege. More precisely, we prove that the pigeonhole principle requires proofs of size $n^{\Omega(k)}$ in parameterized bounded-depth Frege, and, as a special case, in dag-like Parameterized Resolution. This answers an open question posed in [9]. In the opposite direction, we interpret a well-known technique for FPT algorithms as a DPLL procedure for Parameterized Resolution. Its generalization leads to a proof search algorithm for Parameterized Resolution that in particular shows that tree-like Parameterized Resolution allows short refutations of all parameterized contradictions given as bounded-width CNF's.

1 Introduction

Recently, Dantchev, Martin, and Szeider [9] introduced the framework of *parameterized proof complexity*, an extension of the proof complexity approach of Cook and Reckhow to parameterized complexity. One motivation for this is the quest for efficient algorithms which solve optimization problems [10, 11, 16]. Since Resolution is very important for SAT solving, its analogue in this context, Parameterized Resolution, combines these two approaches, and its investigation might provide new insights into proof search for tractable fragments of classically hard problems. Some results in this direction are already outlined in the

* Part of this work was done while the first author was visiting Sapienza University Rome under support of grant N. 20517 by the John Templeton Foundation.

** Supported by grant “Limiti di compressione in combinatoria e complessità computazionale” by Sapienza University Rome.

*** Part of this work was done while the author was at Steklov Mathematical Institute, supported by the Russian Foundation for Basic Research, and at Toyota Technological Institute, Chicago.

work of Gao [13] where he analyzes the effect of the standard DPLL algorithm on the problem of weighted satisfiability for random d -CNF. However, the study of Parameterized Resolution and our understanding of the possible implications for SAT-solving algorithms are still at a very early stage.

More generally, *parameterized complexity* is a branch of complexity theory where problems are analyzed in a different way than in the classical approach: we say that a problem is *fixed-parameter tractable* (FPT) with parameter k if any instance of size n can be solved in time $f(k)n^{O(1)}$ for some computable function f of arbitrary growth. In this setting, classically intractable problems may have efficient solutions for small choices of the parameter, even if the total size of the input is large. Consider e.g. the classical satisfiability problem of finding a truth assignment that satisfies all clauses of a formula in conjunctive normal form. BOUNDED CNF SAT and WEIGHTED CNF SAT are parameterized variants of CNF satisfiability in which the satisfying assignment is required to have Hamming weight at most k or exactly k , respectively. Many parameterized combinatorial problems can be naturally encoded in BOUNDED CNF SAT or WEIGHTED CNF SAT: finding a vertex cover of size at most k , finding a clique of size k , or finding a dominating set of size at most k . In the theory of parameterized complexity, the hardness of both problems is reflected by their W[2]-completeness.

In [9], Dantchev, Martin, and Szeider laid the foundations to study complexity of proofs in a parameterized setting. The dual problem of BOUNDED CNF SAT is that of deciding *parameterized contradictions* PCon: it consists of all pairs (F, k) where F is a propositional formula F which has no satisfying assignment of weight $\leq k$. After considering this notion of propositional *parameterized tautologies*, Dantchev et al. [9] introduced the concepts of parameterized proof systems and of fpt-bounded proof systems (see Section 2 for a discussion). The main motivation behind the work of [9] was that of generalizing the classical approach of Cook and Reckhow to the parameterized case and working towards a separation of parameterized complexity classes as FPT and W[2] by techniques developed in proof complexity. In fact, we obtain an analogous result to the well-known Cook-Reckhow theorem from [8]: A parameterized language L has an fpt-bounded proof system if and only if $L \in \text{para-NP}$ (Theorem 1).

In [9] (tree-like) *Parameterized Resolution* was defined as a refutation system for the set of parameterized contradictions. If $(F, k) \in \text{PCon}$ is defined on variables x_1, \dots, x_n then a (tree-like) *Parameterized Resolution refutation* of (F, k) is a (tree-like) Resolution refutation of $F \cup \{\neg x_{i_1} \vee \dots \vee \neg x_{i_{k+1}} \mid 1 \leq i_1 < \dots < i_{k+1} \leq n\}$. Thus, in (tree-like) Parameterized Resolution we have built-in access to all parameterized clauses of the form $\neg x_{i_1} \vee \dots \vee \neg x_{i_{k+1}}$. All these clauses are available in the system, but when measuring the size of a derivation we only count those which actually appear in the derivation. This concept can be straightforwardly generalized to an arbitrary proof system P , be it dag-like or tree-like, that understands clauses and works with lines.

Dantchev et al. [9] prove an extension of Riis' *gap theorem* [20] and obtain a model theoretic classification for the complexity of tree-like Parameterized Res-

olution refutations for parameterized contradictions originating as propositional encodings of first-order formulas. In particular, their main result implies that tree-like Parameterized Resolution is not fpt-bounded. A similar question for dag-like Parameterized Resolution was left open in [9]. More specifically, they asked if (the parameterized version of) the pigeonhole principle is hard for dag-like Parameterized Resolution.

We answer this question by proving that PHP_n^{n+1} requires proofs of size $n^{\Omega(k)}$ not only in Parameterized Resolution but in the much stronger system of bounded-depth Frege. Our result is in sharp contrast with [9, Proposition 17] that gives efficient proofs of PHP_n^{n+1} in Parameterized Resolution using a more sophisticated encoding with auxiliary variables. We discuss these augmented proof systems in the final Section 5. Our lower bound for the pigeonhole principle is a rather simple application of the method of random restrictions introduced in proof complexity by Haken in his seminal paper [14]. But our choice of parameters is totally different and allows us to kill with the restriction any small prescribed set of parameterized axioms. While the technique is routine, it nonetheless seems to be its first application in the context of *parameterized* complexity, be it computational or proof complexity.

As our second contribution we investigate classes of parameterized contradictions that have short refutations in tree-like Parameterized Resolution. The notion of *kernelization* plays an important role in the theory of parameterized complexity to design fpt-algorithms. Here we propose a notion of *core* for parameterized proof complexity: the core of a parameterized contradiction (F, k) is a subset of clauses $F' \subseteq F$ whose size is bounded by a function of k only, and such that (F', k) is still a parameterized contradiction. We observe that if a formula has a core, then it can be efficiently refuted in tree-like Parameterized Resolution with a refutation of size independent of the size of F . As an immediate consequence, several examples of formulas hard for tree-like Resolution are instead efficiently refutable in the parameterized case: pebbling contradictions, linear ordering principles, graph pigeonhole principles, and colorability principles. But sometimes a core of a formula is not explicit or immediate to find. In Theorem 3 we prove that contradictions of bounded width have a core and thus very efficient tree-like Parameterized Resolution refutations.

2 Parameterized Proof Complexity

A *parameterized language* is a language $L \subseteq \Sigma^* \times \mathbb{N}$. For an instance (x, k) , we call k the *parameter of* (x, k) . A parameterized language L is *fixed-parameter tractable* if L has a deterministic decision algorithm running in time $f(k)|x|^{O(1)}$ for some computable function f . The class of all fixed-parameter tractable languages is denoted by FPT.

Besides FPT there is a wealth of complexity classes containing problems which are not believed to be fixed-parameter tractable. The most prominent classes lie in the *weft hierarchy* forming a chain $\text{FPT} \subseteq \text{W}[1] \subseteq \text{W}[2] \subseteq \dots \subseteq \text{W[P]}$. All these classes are included in **para-NP** (see Definition 6). The classes of the weft hierarchy are usually defined as the closure of a canonical problem

under fpt-reductions. For $W[2]$ this canonical problem is WEIGHTED CNF SAT containing instances (F, k) with a propositional formula F in CNF and a parameter $k \in \mathbb{N}$. WEIGHTED CNF SAT asks whether F has a satisfying assignment of weight exactly k , where the weight of an assignment α , denoted as $w(\alpha)$, is the number of variables that α assigns to 1. Instead of asking for an assignment α with $w(\alpha) = k$ we can also ask for α with $w(\alpha) \leq k$ and still get the $W[2]$ -complete problem BOUNDED CNF SAT (cf. [9] or the full version of this paper for the proof of its $W[2]$ -completeness).

Like in the classical duality between tautologies and satisfiability, the complement of BOUNDED CNF SAT is a complete problem for $\text{co}W[2]$:

Definition 1 (Dantchev, Martin, Szeider [9]). A parameterized contradiction is a pair (F, k) consisting of a propositional formula F , given as a CNF, and $k \in \mathbb{N}$ such that F has no satisfying assignment of weight $\leq k$. We denote the set of all parameterized contradictions by PCon .

Next we discuss the general definition of a parameterized proof system from [9].

Definition 2 (Dantchev, Martin, Szeider [9]). A parameterized proof system for a parameterized language $L \subseteq \Sigma^* \times \mathbb{N}$ is a function $P : \Sigma^* \times \mathbb{N} \rightarrow \Sigma^* \times \mathbb{N}$ such that $\text{rng}(P) = L$ and $P(x, k)$ can be computed in time $f(k)|x|^{O(1)}$ with some computable function f .

The purpose of the second argument in P remains a little bit unclear to us since all natural proof systems we can think of do not have this feature. Thus, we propose the following simplification.

Definition 3. A proof system for a parameterized language $L \subseteq \Sigma^* \times \mathbb{N}$ is a polynomial-time computable function $P : \Sigma^* \rightarrow \Sigma^* \times \mathbb{N}$ such that $\text{rng}(P) = L$.

Now we would like to show that both versions are even formally equivalent in the sense that a parameterized language has a proof system in which all strings possess “short” proofs if and only if it has a parameterized proof system with this property. First we have to formalize the notion of “short”. In the framework of [9] it goes as follows:

Definition 4 (Dantchev, Martin, Szeider [9]). A parameterized proof system P for a parameterized language L is fpt-bounded if there exist computable functions f and g such that every $(x, k) \in L$ has a P -proof (y, k') with $|y| \leq f(k)|x|^{O(1)}$ and $k' \leq g(k)$.

Again, our analogue is simpler.

Definition 5. A proof system P for a parameterized language L is fpt-bounded if there exists a computable function f such that every $(x, k) \in L$ has a P -proof of size at most $f(k)|x|^{O(1)}$.

Recall that by the theorem of Cook and Reckhow [8], the class of all languages with polynomially bounded proof systems coincides with NP . To obtain a similar result in the parameterized world, we use the following parameterized version of NP .

Definition 6 (Flum, Grohe [11]). *The class para-NP contains all parameterized languages which can be decided by a nondeterministic Turing machine in time $f(k)|x|^{O(1)}$ for a computable function f .*

Theorem 1. *Let $L \subseteq \Sigma^* \times \mathbb{N}$ be a parameterized language. Then the following statements are equivalent: (1) There exists an fpt-bounded proof system for L . (2) There exists an fpt-bounded parameterized proof system for L . (3) $L \in \text{para-NP}$.*

Proof (Sketch). (1) equivalent to (2): by padding the proof, the proof verification procedure can be made polynomial in proof length. (1) equivalent to (3): a proof can be guessed nondeterministically and then verified. \square

2.1 Parameterized Versions of Ordinary Proof Systems

A *literal* is a positive or negated propositional variable and a *clause* is a set of literals. The *width* of a clause is the number of its literals. A clause is interpreted as the disjunction of its literals and a set of clauses as the conjunction of the clauses. Hence clause sets correspond to formulas in CNF.

The system of *Parameterized Resolution* was introduced by Dantchev, Martin, and Szeider [9]. Parameterized Resolution is a refutation system for the set PCon of parameterized contradictions (cf. Definition 1). Given a set of clauses F in variables x_1, \dots, x_n with $(F, k) \in \text{PCon}$, a *Parameterized Resolution refutation* of (F, k) is a Resolution refutation of

$$F \cup \{\neg x_{i_1} \vee \dots \vee \neg x_{i_{k+1}} \mid 1 \leq i_1 < \dots < i_{k+1} \leq n\}. \quad (1)$$

Thus, in Parameterized Resolution we have built-in access to all parameterized clauses of the form $\neg x_{i_1} \vee \dots \vee \neg x_{i_{k+1}}$. All these clauses are available in the system, but when measuring the size of a derivation we only count those which appear in the derivation. Note that Parameterized Resolution is actually a proof system for PCon in the sense of Definition 3, i.e., verification proceeds in polynomial time. This definition allows the following straightforward generalization.

Definition 7. *Let $P : \Sigma^* \rightarrow \text{Con}$ be an ordinary proof system for the language Con of all (ordinary) CNF contradictions. We define the parameterized version \hat{P} of P by letting $\hat{P}(F, k, x) = (F, k)$ whenever $P(x)$ is an arbitrary subset of the set of axioms (1). If $P(x)$ does not have this form, $\hat{P}(F, k, x)$ outputs something trivial.*

The only specific proof system we would like to comment on is tree-like Parameterized Resolution (as it will be needed in Section 4). As explained in [9], a tree-like Parameterized refutation of (F, k) can equivalently be described as a *boolean decision tree*. A boolean decision tree for (F, k) is a binary tree where inner nodes are labeled with variables from F and leafs are labeled with clauses from F or parameterized clauses $\neg x_{i_1} \vee \dots \vee \neg x_{i_{k+1}}$. Each path in the tree corresponds to a partial assignment where a variable x gets value 0 or 1 according to whether the path branches left or right at the node labeled with x . The

condition on the decision tree is that each path α must lead to a clause which is falsified by the assignment corresponding to α . Therefore, a boolean decision tree solves the *search problem* for (F, k) which, given an assignment α , asks for a clause falsified by α . It is easy to verify that each tree-like Parameterized Resolution refutation of (F, k) yields a boolean decision tree for (F, k) and vice versa, where the size of the Resolution proof equals the number of nodes in the decision tree.

An embarrassing fact about Parameterized Proof Complexity (brought to our attention by an anonymous referee of a previous version of this paper) is that, as defined in Definition 7, \widehat{P} is *never* bounded for some dull reasons.

Example 1. Let (F, k) be the parameterized contradiction in which F is the set of positive clauses $\{x_{1,1} \vee \dots \vee x_{1,n}, \dots, x_{k+1,1} \vee \dots \vee x_{k+1,n}\}$. Then in order to make this set even *semantically* invalid, one has to append to it all n^{k+1} parameterized axioms of the form $\neg x_{1,j_1} \vee \dots \vee \neg x_{k+1,j_{k+1}}$.

Obviously, this is not the kind of phenomena we want to study (and not the kind of methods we want to develop) so we have to try to somehow isolate such pathological examples. One approach (borrowed from circuit complexity) would be simply to declare some parameterized contradictions “natural”, “interesting” or “explicit” without giving precise definitions or even revealing exact reasons for this classification. Another possibility (that we adopt in this paper) is to *formally* restrict the set of contradictions we are interested in.

Definition 8. A parameterized contradiction (F, k) is *strong* if F itself is a contradiction. A proof system P for the set PCon is *weakly fpt-bounded* if there exists a computable function f such that every **strong** $(F, k) \in \text{PCon}$ has a P -proof of size at most $f(k)|F|^{O(1)}$.

One reason to introduce this restriction is that “interesting” contradictions are almost always strong. In fact, the only exception we are aware of (even if it is the one that inspired almost all material in Section 4) is the vertex cover problem.

On a more philosophical level, the concept of a strong parameterized contradiction intends to capture the idea that the new knowledge provided by parameterized axioms should be rather thought of as a *helper* or an *additional feature* made available to already existing DPLL algorithms rather than being the prime source of the *validity* of the statement. It is perhaps useful to note at this point that there is no monotonicity on k in Definition 8, either way. Since, as k grows, our size constraint $f(k)|F|^{O(1)}$ gets looser, but, on the other hand, the helper axioms become less helpful.

Finally, we are not aware of any analogue of Example 1 for strong parameterized contradictions.

Yet another possibility to get rid of this example is to try to encode parameterized axioms in (1) in a more economical way (so that their number stays small), possibly using some auxiliary variables. For Parameterized Resolution this possibility was discussed already in [9], and we continue this discussion in a broader context in Section 5.

3 Parameterized Bounded-Depth Frege is Not Weakly fpt-Bounded

The *pigeonhole principle* PHP_n^{n+1} uses variables $x_{i,j}$ with $i \in [n+1]$ and $j \in [n]$, indicating that pigeon i goes into hole j . PHP_n^{n+1} consists of the clauses $\bigvee_{j \in [n]} x_{i,j}$ for all pigeons $i \in [n+1]$ and $\neg x_{i_1,j} \vee \neg x_{i_2,j}$ for all choices of two distinct pigeons $i_1, i_2 \in [n+1]$ and a hole $j \in [n]$.

Let F_d be the fragment of the Frege system over de Morgan basis $\{\neg, \wedge, \vee\}$ that operates with formulas of logical depth at most d .

Theorem 2. *For any fixed $d, k \geq 0$ and all sufficiently large n , any refutation of (PHP_n^{n+1}, k) in \widehat{F}_d , the parameterized version of F_d , requires size $\geq n^{k/5}$.*

Note that, somewhat surprisingly, d does not appear in the final bound at all (although it implicitly appears in the bound assumed in the “sufficiently large” premise).

Proof. Choose uniformly at random a set I of $n - \sqrt{n}$ pigeons and match them with a set J of $n - \sqrt{n}$ uniformly chosen holes. Such partial matching f induces the following natural partial assignment of the variables of PHP_n^{n+1} :

$$\begin{aligned} x_{i,j} &= 1 && \text{whenever } i \in I \text{ and } f(i) = j, \\ x_{i,j} &= 0 && \text{whenever } i \in I \text{ and } f(i) \neq j, \\ x_{i,j} &= 0 && \text{whenever } j \in J \text{ and there exist } i' \neq i \text{ such that } f(i') = j, \text{ and} \\ x_{i,j} &= \star && \text{otherwise.} \end{aligned}$$

We claim that with non-zero probability such partial assignment satisfies all parameterized axioms used in the refutation, as long as there are at most $n^{k/5}$ of them. (Notice that we do not care if such assignment falsifies unused parameterized axioms.) Before proving this claim, we show how the theorem follows.

The refutation, restricted with such assignment, does not contain parameterized axiom anymore. Thus it is a classical F_d -refutation for the restricted formula, which in turn is equivalent (up to a re-indexing of pigeons and holes) to $PHP_{\sqrt{n}}^{\sqrt{n}+1}$. Such refutation must be of size at least $2^{n^{c_d}}$ [17, 15] for some $c_d > 0$, thus bigger than $n^{k/5}$ if n is sufficiently large. This concludes the proof.

The missing part is to show that the probabilistic choice of the partial matching realizes the desired properties with positive probability. Consider a parameterized axiom $\neg x_{i_1, j_1} \vee \dots \vee \neg x_{i_{k+1}, j_{k+1}}$. If there are two equal indexes j_a and j_b for $a \neq b$, then such axiom is just a weakening of a standard clause of the pigeonhole principle and does not need any special treatment.

We can now focus on a parameterized axiom in which exactly $k+1$ holes are represented: the probability that such axiom fails to be satisfied is the probability that all x_{i_l, j_l} are either true or unassigned for $1 \leq l \leq k+1$. Let $J_0 = \{j_1, \dots, j_{k+1}\}$ be the set of all holes represented in our axiom. The probability that the support J of our random restriction contains at most $k/2$ of them (and hence the complement to J that has size \sqrt{n} contains at least $k/2$ of them) is bounded by $\binom{k+1}{k/2} \cdot \frac{\binom{n-k/2}{\sqrt{n}-k/2}}{\binom{n}{\sqrt{n}}} \leq 2^{k+1} n^{-k/4}$. And, conditioned by

the event $|J \cap J_0| \geq k/2$, the probability that every hole $j_a \in J \cap J_0$ is sent by the matching f to the right pigeon i_a (so that x_{i_a, j_a} is not set to 0) is at most $(n - k/2)^{-k/2}$. Thus, the overall probability that our random partial assignment does not satisfy an individual parameterized axiom is bounded by $2^{k+1}n^{-k/4} + (n - k/2)^{-k/2} < n^{-k/5}$ for sufficiently large n . By the union bound, if our refutation has size $\leq n^{k/5}$, then for at least one particular choice of f the corresponding assignment satisfies all parameterized axioms actually used in the refutation. As we already observed, this concludes the proof. \square

The same proof works for weaker versions of the pigeonhole principle, like functional or onto, and it works up to $k \leq n^{\epsilon_d}$, where $\epsilon_d > 0$ depends only on the depth of the Frege system. If we consider Parameterized Resolution instead of parameterized bounded-depth Frege, our proof applies also to the pigeonhole principle with arbitrarily many pigeons. See the full version of this paper for details.

4 Cores and Small Refutations

The notion of *kernelization* plays an important role in the theory of parameterized complexity. A kernelization for a parameterized language L is a polynomial-time procedure $A : \Sigma^* \times \mathbb{N} \rightarrow \Sigma^* \times \mathbb{N}$ such that for each (x, k) it holds that (1) $(x, k) \in L$ if and only if $A(x, k) \in L$ and (2) if $A(x, k) = (x', k')$, then $k' \leq k$ and $|x'| \leq f(k)$ for some computable function f independent of $|x|$.

It is clear that if a parameterized language admits a kernelization then it is fixed-parameter tractable. The converse is also true for decidable languages (cf. [12]). For parameterized proof complexity we suggest a similar notion of core for parameterized contradictions:

Definition 9. A core for a set $\Gamma \subseteq \text{PCon}$ of parameterized contradictions is a mapping which maps every $(F, k) \in \Gamma$ to a subset $F' \subseteq F$ of clauses satisfying the following conditions: (1) F' contains at most $f(k)$ variables and (2) (F', k) is a parameterized contradiction, where f is a computable function depending only on the mapping.

We will sometimes abuse terminology by saying that a set of clauses $F' \subseteq F$ is a core of F when it is clear from the context that F is a member of a family of parameterized contradictions and that F' can be chosen for any k and any F in that family.

Clearly, any $k + 1$ positive clauses of width $f(k)$ and with pairwise disjoint sets of variables make a core that we will call a *trivial core*. It is very easy to come up with many parameterized contradictions (pebbling contradictions, colorability, sparse pigeonhole principle etc.) that possess trivial cores. Thus an interesting question is the following: do there exist “natural” parameterized contradictions that possess only non-trivial cores? And do we have a “parameterized automatizability”, i.e., is it easy to find a core once we know that it exists?

Our motivating example for the latter question is the *vertex cover* problem. A vertex cover for a graph G is a set $C \subseteq V(G)$ such that for any $\{u, v\} \in E(G)$

either $u \in C$ or $v \in C$ or both. To determine whether G has a vertex cover of size at most k there is a well-known [10, Chapter 3] fixed parameter tractable algorithm (here the parameter is k). This algorithm is based on the following observation: if a vertex is not in C , then all its neighbors must be in C . The algorithm is a simple recursive procedure which focuses on an arbitrary vertex u , and on its neighbors v_1, \dots, v_l : if neither $G \setminus \{u\}$ has a vertex cover of size $k-1$ nor $G \setminus \{u, v_1, \dots, v_l\}$ has a vertex cover of size $k-l$, then G has no vertex cover of size k .

This is interpretable as a parameterized DPLL procedure on the 2-CNF $F_G = \bigwedge_{\{u,v\} \in E(G)} (x_u \vee x_v)$ where x_u indicates whether $u \in C$. The DPLL procedure fixes an arbitrary variable x_u and branches on it. When $x_u = 1$, then the DPLL algorithm proceeds with analyzing $F_G \upharpoonright_{x_u=1}$ which is equal to $F_{G \setminus \{u\}}$. When $x_u = 0$, then $x_{v_1} = 1, \dots, x_{v_l} = 1$ by unit propagation. Thus the DPLL proceeds on formula $F_G \upharpoonright_{\{x_u=0, x_{v_1}=1, \dots, x_{v_l}=1\}} = F_{G \setminus \{u, v_1, \dots, v_l\}}$. If at any point the DPLL has more than k variables set to one, it stops and backtracks.

And now we establish a far-reaching generalization of this example.

Theorem 3. *If F is a CNF of width d and (F, k) is a parameterized contradiction, then (F, k) has a tree-like Parameterized Resolution refutation of size $O(d^{k+1})$. Moreover, there is an algorithm that for any (F, k) either finds such tree-like refutation or finds a satisfying assignment for F of weight $\leq k$. The algorithm runs in time $O(|F| \cdot k \cdot d^{k+1})$.*

A related result was obtained in [7, Theorem 12]. Notice that while BOUNDED CNF SAT and WEIGHTED CNF SAT are both W[2]-complete, BOUNDED d -CNF SAT is in FPT and WEIGHTED d -CNF SAT is known to be W[1]-complete. This means that reducing the case of exact weight to bounded weight requires large clauses unless $\text{FPT} = \text{W}[1]$. We state two interesting consequences of Theorem 3.

Corollary 1. *For each $d \in \mathbb{N}$, the set of all parameterized contradictions in d -CNF has a core.*

Proof. The refutations constructed in Theorem 3 contain $O(d^k)$ initial clauses in $O(d^{k+1})$ variables. These clauses form a core. \square

The following corollary expresses some restricted form of automatizability (cf. also the discussion in Section 5).

Corollary 2. *If $\Gamma \subseteq \text{PCon}$ has a core, then there exists an fpt-algorithm which on input $(F, k) \in \Gamma$ returns both a core and a refutation of (F, k) .*

Proof. Let Γ have a core of size $f(k)$. Then the core only contains clauses of width $\leq f(k)$. On input (F, k) we run the algorithm of Theorem 3 on the CNF formula consisting of all clauses of F with width $\leq f(k)$. This yields a core together with its refutation. \square

Finally, we comment on the question whether the existence of a core is a necessary condition for a parameterized contradiction to have an fpt-bounded

refutation in tree-like Parameterized Resolution. A trivial counterexample to this conjecture is made by the CNF $(x_1 \vee x_2 \vee \dots \vee x_n) \wedge \neg x_1 \wedge \dots \wedge \neg x_n$. A more interesting example—a version of the linear ordering principle—will be included in the full version of this paper.

5 Discussion and Open Problems

Is the parameterized proof system \widehat{P} from Definition 7 the most natural way to define the parameterized analogue of P ? The answer depends on the original proof system P , of course. The main (unspoken) reason why [9] defined it in this way is simply because weak proof systems cannot directly talk about the weight of the input. Let us first discuss two familiar systems that are strong enough to overcome this limitation: Frege and Cutting Planes.

The problem of getting super-polynomial lower bounds for the Frege proof system F is one of the biggest open problems in Logic and Theoretical Computer Science. Lower bounds for its parameterized version \widehat{F} seem even harder to achieve for strong contradictions (as we just add new axioms). A similar conclusion remains true if we combine all parameterized axioms into one (using e.g. [6]) but allow arbitrary parameterized contradictions, not necessarily strong.

The case of Cutting Planes (CP) is way more interesting. First of all, we do not seem to know lower bounds even for the “canonical” version \widehat{CP} :

Question 1. Is \widehat{CP} weakly fpt-bounded?

This, of course, is yet another reflection of the mysterious status of this proof system: the only known lower bounds for it are based on very indirect methods (interpolation, see [5, 18]) and no direct, combinatorial proof is currently known. And if we try to generalize the methods from [5, 18] (at least in a straightforward way) then we immediately arrive at a problem in parameterized circuit complexity that seems to be widely open (at least, we do not see how known methods can be applied to it).

Question 2. Find an explicit *partial* monotone function (a.k.a. a monotone promise problem) $f : [n]^{\leq k} \rightarrow \{0, 1\}$ defined only on inputs of Hamming weight $\leq k$ that does not possess monotone circuits of size $f(k)n^{O(1)}$.

The problem of finding (say) a \sqrt{k} -clique does become easy in this context.

For weaker proof systems, [9, Section 4] proposed to use auxiliary variables. Their suggestion was to add new “pigeonhole variables” $p_{i,j}$ ($i \in [n]$, $j \in [k]$) and “pigeonhole clauses” $\neg x_i \vee \bigvee_{j \in [k]} p_{i,j}$ for all $i \in [n]$ (pigeon clauses) and $\neg p_{i_1,j} \vee \neg p_{i_2,j}$ for all $i_1 \neq i_2 \in [n]$, $j \in [k]$ (hole clauses), where x_1, \dots, x_n are the original variables. Remarkably, they proved that the pigeonhole principle has fpt-bounded refutations in this version of Parameterized Resolution.

The disturbing Example 1 turns into an instance of PHP_k^{k+1} with large “metapigeons” that has an fpt-bounded proof (e.g., the straightforward adaption of the rectangular proof from [19, Example 1]). Thus, following [9], we ask:

Question 3. Is Parameterized Resolution with auxiliary variables fpt-bounded?

Let us now point out that there is an interesting and well-studied class of contradictions for which the difference between these two encodings disappears, and these are independent set principles. Following [2], let G be a graph $[n]$ in which vertices are split into k subsets V_1, \dots, V_k of size n/k each called *blocks*. The principle $\alpha_{block}(G, k)$ encodes the fact that G has a block-respecting independent set of size k ; it has the variables x_v ($v \in [n]$) and the axioms $\neg x_u \vee \neg x_v$ for all $(u, v) \in E(G)$ (edge clauses), $\bigvee_{v \in V_i} x_v$ for all $i \in [k]$ (block clauses), and $\neg x_u \vee \neg x_v$ for all $u \neq v$ in the same block (1–1 clauses).

The fact that all satisfying assignments have at most k ones is already built in this principle: all parameterized axioms are subsumed by the 1–1 clauses above. Auxiliary clauses in the sense of [9] (both pigeon and holes) also do not help to reduce the refutation size, as witnessed by the substitution of the pigeonhole variables $p_{v,j} \mapsto 0$ if $v \notin V_j$ and $p_{v,j} \mapsto x_v$ if $v \in V_j$. Thus, we are also asking the following specific form of Question 3:

Question 4. Do the principles $\alpha_{block}(G, k)$ always have fpt-bounded Resolution refutations as long as the graph G does not contain block-respecting independent sets of size k ?

One good candidate for a lower bound here would be Erdős-Rényi random graphs $G(n, p)$ for an appropriately chosen value of p . Such lower bound has been recently proved for tree-like Resolution in [4].

Let us recall that a proof system P is *automatizable* if there exists an algorithm which for a tautology F with a P -proof of size S finds a P -proof for F of size at most $S^{O(1)}$ and runs in time $S^{O(1)}$. Alekhovich and Razborov [1] proved that if (classical) Resolution or tree-like Resolution were automatizable, then $W[P]$ would coincide with FPR, the randomized version of FPT. On the other hand, tree-like Resolution is quasi-polynomially automatizable (see e.g. [3]).

We point out that the concept of quasi-polynomial automatizability is meaningless in the context of Parameterized Resolution, because every $(F, k) \in \text{PCon}$ with $|F| = n$ has a refutation of size $c \cdot \binom{n}{k+1}$ for some constant c . If $k \leq \log n$ this is smaller than $n^{\log n}$; otherwise $\binom{n}{k+1} \leq 2^{(k+1)^2}$ which is fpt with respect to k . On the contrary, the concept of polynomial automatizability can be extended to parameterized proof systems in an obvious way. Thus, we ask:

Question 5. Is (tree-like) Parameterized Resolution, with or without auxiliary variables, fpt-automatizable or fpt-automatizable w.r.t. strong contradictions? That is, does there exist an algorithm that for any (strong) parameterized contradiction $(F, k) \in \text{PCon}$ outputs its refutation within time $f(k)S^{O(1)}$, where S is the minimal possible size of a parameterized refutation of (F, k) ?

Naturally, unconditional results of this sort are completely out of reach for the moment, so we are willing to allow here any reasonable complexity assumption (that will most likely reside in the realm of parameterized complexity itself).

References

1. M. Alekhovich and A. A. Razborov. Resolution is not automatizable unless $W[P]$ is tractable. *SIAM Journal on Computing*, 38(4):1347–1363, 2008.
2. P. Beame, R. Impagliazzo, and A. Sabharwal. The resolution complexity of independent sets and vertex covers in random graphs. *Computational Complexity*, 16(3):245–297, 2007.
3. P. Beame, R. M. Karp, T. Pitassi, and M. E. Saks. The efficiency of resolution and Davis–Putnam procedures. *SIAM J. Comput.*, 31(4):1048–1075, 2002.
4. O. Beyersdorff, N. Galesi, and M. Lauria. Parameterized complexity of DPLL search procedures. In *Proc. 14th International Conference on Theory and Applications of Satisfiability Testing*, 2011. To appear.
5. M. L. Bonnet, T. Pitassi, and R. Raz. Lower bounds for cutting planes proofs with small coefficients. *The Journal of Symbolic Logic*, 62(3):708–728, 1997.
6. S. R. Buss. Polynomial size proofs of the propositional pigeonhole principle. *The Journal of Symbolic Logic*, 52:916–927, 1987.
7. Y. Chen and J. Flum. The parameterized complexity of maximality and minimality problems. *Annals of Pure and Applied Logic*, 151(1):22–61, 2008.
8. S. A. Cook and R. A. Reckhow. The relative efficiency of propositional proof systems. *The Journal of Symbolic Logic*, 44(1):36–50, 1979.
9. S. S. Dantchev, B. Martin, and S. Szeider. Parameterized proof complexity. In *Proc. 48th IEEE Symposium on the Foundations of Computer Science*, pages 150–160, 2007.
10. R. G. Downey and M. R. Fellows. *Parameterized Complexity*. Springer-Verlag, Berlin Heidelberg, 1999.
11. J. Flum and M. Grohe. Describing parameterized complexity classes. *Information and Computation*, 187(2):291–319, 2003.
12. J. Flum and M. Grohe. *Parameterized Complexity Theory*. Springer-Verlag, Berlin Heidelberg, 2006.
13. Y. Gao. Data reductions, fixed parameter tractability, and random weighted d-CNF satisfiability. *Artificial Intelligence*, 173(14):1343–1366, 2009.
14. A. Haken. The intractability of resolution. *Theor. Comput. Sci.*, 39:297–308, 1985.
15. J. Krajíček, P. Pudlák, and A. Woods. Exponential lower bounds to the size of bounded depth Frege proofs of the pigeonhole principle. *Random Structures and Algorithms*, 7(1):15–39, 1995.
16. R. Niedermeier. *Invitation to Fixed-Parameter Algorithms*. Oxford Lecture Series in Mathematics and Its Applications. Oxford University Press, 2006.
17. T. Pitassi, P. Beame, and R. Impagliazzo. Exponential lower bounds for the pigeonhole principle. *Computational Complexity*, 3:97–140, 1993.
18. P. Pudlák. Lower bounds for resolution and cutting planes proofs and monotone computations. *The Journal of Symbolic Logic*, 62(3):981–998, 1997.
19. A. Razborov, A. Wigderson, and A. Yao. Read-once branching programs, rectangular proofs of the pigeonhole principle and the transversal calculus. *Combinatorica*, 22(4):555–574, 2002.
20. S. Riis. A complexity gap for tree resolution. *Computational Complexity*, 10(3):179–209, 2001.